

How a model sees the world

Every answer is shaped by what's in the context: choose it, measure it, pay for it.

"Context is the model's working reality."

Ask on every project

1. What is in the context on a typical call, and how many tokens does each part take?
2. Which parts of that context does this answer actually need?
3. Where do the key rule and facts sit: at the start, buried mid-way, or by the question?
4. What gets trimmed when the window fills, and would anyone notice?
5. Do any sources in the context contradict each other or include outdated versions?

Rules of thumb

- Count tokens with the model's own tokenizer; don't estimate from word counts.
- Send only what this request needs; cut stale versions and unrelated history.
- Keep critical rules in instructions sent every call; restate them near the question.
- Estimate tokens per call × calls per day before deciding how much context to send.

Red flags

- ⚠ "Just put the whole handbook in the prompt, the window is big enough"
- ⚠ Old and current versions of a policy side by side in the same context
- ⚠ Nobody knows the token count or cost of a typical request
- ⚠ Testing only on short chats, never on long conversations

Key terms

Token	A piece of text (word, word part or symbol) that the model reads as one number.
Embedding	A long list of numbers representing a token; related meanings sit close together.
Attention	How each token weighs earlier tokens to pull in what's relevant to it.
Context window	The hard token limit for everything in one call, the answer included.
Dilution	Relevant text used less reliably because unrelated text competes with it.

CHECKLIST

Context budget check

- | | |
|---|--|
| ☐ List every source in the context: instructions, history, documents, tool output, question. | ☐ Count tokens per source with the real tokenizer and note the biggest share. |
| ☐ Cut what this request doesn't need: stale versions, unrelated turns, noisy tool output. | ☐ Keep key rules in instructions, label key facts, and put the question near the end. |
| ☐ Check what gets trimmed near the window limit, and leave room for the answer. | ☐ Multiply tokens per call by calls per day, and time a long request end to end. |