

Agents without the mythology

Autonomy is a cost: give the model a decision only where code can't cover the cases.

"An agent is a loop with model-selected decisions."

Ask on every project

1. Who decides the next step here: our code, or the model?
2. Are the steps the same for every input, or do they depend on what it finds?
3. What stops the loop: a step limit, a budget, a timeout or a human?
4. If step 3 is wrong, which later step or person notices it, and when?
5. If a run fails halfway, what state do we need to resume it or undo it?

Rules of thumb

- Make known steps fixed code; keep model decisions for the parts that truly vary.
- Set step, budget and time limits in code, and decide what happens when one is hit.
- Check key results before they enter state; later steps will treat them as facts.
- Put approval on few, high-risk actions, and show the approver the evidence.

Red flags

- ⚠ It's called an agent, but nobody can say what the model decides
- ⚠ No step limit or budget, so a stuck run keeps retrying the same action
- ⚠ One person approves hundreds of actions a day without reading them
- ⚠ More agents are added before a simple workflow was tried and measured

Key terms

Workflow	Code fixes the order of steps; the model only fills in some of them.
Agent loop	The model picks the next action, the system runs it, and the result feeds the next turn.
Router	One model decision that sends an input down one of several fixed paths.
State	What a run carries forward: goal, history, notes and changes made in other systems.
Blast radius	How much damage one wrong action can do before anyone can stop or undo it.

DECISION GUIDE

How much autonomy?

If the steps are the same for every input	→ fixed workflow; the model only fills in steps
If inputs fall into a few known paths	→ router: one model decision, then fixed code
If the path depends on what it finds and tools are read-only	→ bounded agent with caps
If an action sends, spends, deletes or changes production	→ a human approves it first
If an action can't be undone and approval isn't practical	→ don't expose it as a tool
If you can't yet measure whole runs	→ stay one level less autonomous until you can