

Building AI you can trust

If you can't name the number that would make you roll back, you don't know it works.

"A demo is not a system."

Ask on every project

1. Which written cases define "working" for this feature, and who agreed on them?
2. Which quality checks can plain code do before we reach for a model judge?
3. Which failure type is most common, and what does each type cost us?
4. What reruns the set when a prompt, model, index or tool changes?
5. What p95 latency and cost per request can this use case afford?

Rules of thumb

- Grade with code first, a calibrated judge next, and people to set the rubric.
- Turn every production failure into a golden-set case before you fix it.
- Compare pass rates over repeated runs, not single outputs, before judging a change.
- Block a release when golden-set scores drop below the stored baseline.

Red flags

- ⚠ Quality is shown only with hand-picked demo examples
- ⚠ A single accuracy number with no breakdown by failure type
- ⚠ A model judge in use with no check against human labels
- ⚠ The app calls an unpinned model name and nothing reruns the set

Key terms

Golden set	A small, curated, versioned set of real cases the team agrees defines "working".
Grader	Anything that scores an output: a code check, a human rater or a model judge.
LLM-as-a-judge	A model that grades outputs against a written rubric; calibrate it against people.
Regression test	Rerunning the set after a change to catch cases that used to pass and now fail.
p95 latency	The response time 95% of requests beat; it shows what the slow cases feel like.

TEMPLATE

One-page evaluation plan

Scope: the feature, its users, and what a costly error looks like

Golden set: sources (logs, tickets, past failures), size, hard cases, who maintains it

Graders: one per criterion (code check, calibrated judge or human), with the rubric

Pass thresholds: per criterion and per failure type, plus the baseline to compare against

Regression trigger: which changes rerun the set, and what score drop blocks a release

Budgets and owner: cost per request, p95 latency, and who acts when a number slips